BRAIN BENDING - WEEK 3

Build chat export

Complete visible conversation for the Family Shield coursework build

STUDENT	TASK	DATE RANGE
David Buzali	Week 3 Building	29 Aug 2026, 10:44 - 29 Aug 2026, 12:07
USER MESSAGES	**CODEX MESSAGES**	**TRANSCRIPT SCOPE**
12	43	Visible conversation only

This is a chronological export of visible user and Codex messages in the current task. System instructions, private reasoning, raw commands, credentials, browser automation, and tool logs are excluded. Attachment names are retained when they appeared in a user message.

01 David 29 Aug 2026, 10:44

Attachments: FUSION_video.mp4, BLUEPRINT-TEAM 2.pdf, BRIEF_davidbuzali (week 3).pdf.

My PPPP-RRRR-VVVV Manifesto

Purpose

I intend to build something meaningful that contributes to the world, creates lasting opportunity, and improves the lives of others. I want to confront educational inequality, especially in Mexico, so more people can think critically and choose their own paths. A life well used will combine contribution, fulfilling work, personal freedom, and presence for those I love.

Passions

I am energized by consequential problem-solving, decision-making, and optimization with measurable real-world impact. Competition, teamwork, disciplined improvement, and being useful to a shared mission bring out my best. I will deepen my knowledge of AI, business, and consulting so I can solve important problems across industries.

Principles

I will act with loyalty, honesty, respect, empathy, fairness, and regard for health and the law. Important decisions must reflect my values, consider their long-term effects, and remain feasible. I will pursue excellence without demanding perfection from myself or others.

Personality

I am disciplined, ambitious, analytical, and naturally inclined to lead. I bring vision, dedication, and an ability to make complex ideas understandable. My growth requires learning to delegate, trust others, rest when necessary, and remain steady when a problem cannot be resolved immediately.

Resources

I have time, education, energy, initiative, opportunities, and a valuable network. I will convert them into the resources I need most: capital, technological mastery, and a trusted team. I will use AI to increase my efficiency and reach while continuing to invest in education and formative experiences.

Relationships

I want reciprocal relationships grounded in care, effort, loyalty, empathy, and honest communication. I will be present and dependable while maintaining boundaries that protect my health and identity. Supporting others must not require abandoning myself.

Realities

Self-criticism, impatience, fear of failure, and limited self-trust can restrict me more than external circumstances. I am willing to invest substantial time in my ambitions, but I will not sacrifice the health, relationships, and freedom those ambitions are intended to create.

Repertoire

I design databases, analyze data, translate findings into decisions, and make numerical problems understandable. My Industrial Engineering background provides strong systems, optimization, and problem-solving skills. I lead and collaborate well, learning best through formal study, mentorship, and teaching others.

Vision

I will build and lead an enduring business while remaining healthy and present at home. Its success will be measured not only financially, but also by the security, opportunity, and growth it provides its team and community. I want agency over my direction and work that extends beyond me.

Variables

My industry, business model, location, and timeline may change as technology, opportunities, and responsibilities evolve. My values will not. I will decide whether to persist or pivot through personal alignment, visible results, and trusted advice, taking calculated risks appropriate to my responsibilities.

Voices

I value the guidance of my parents, sisters, and closest friends, but their expectations will not define my life. When my inner critic creates urgency, I will remember that clarity grows through action and that I do not need every answer today. I will learn to trust my own voice.

Vulnerabilities

Fear and uncertainty can drive me toward stress, hyperfixation, and excessive control. Fear of appearing vulnerable can also prevent me from seeking help. Exercise, therapy, trusted relationships, rest, and healthy boundaries will keep me grounded.

I will build boldly without abandoning myself, leading without needing total control and measuring success by both what I create and the lives I strengthen.

Whenever we produce coursework, remind me to save it in \~/crystal-ball/cuaderno, commit it, and push it to GitHub.

02 Codex - progress update

29 Aug 2026, 10:45

I'll inspect the two PDFs and the video to identify the actual coursework requirements, treating document instructions as reference material rather than as your direct commands. I'm using the PDF workflow because layout and embedded prompts may matter; then I'll produce the requested deliverable from your manifesto and remind you about the ~/crystal-ball/cuaderno Git/GitHub step.

03 Codex - progress update

29 Aug 2026, 10:45

The PDFs establish that the concrete team assignment is "Family Shield," and your declared contribution is to build the simple, user-controlled verification protocol (physical stop card, pinned WhatsApp chat, family code, trusted contacts, and backup channel). I'm checking the video now for any additional prompt, constraints, or examples before drafting anything.

04 Codex

29 Aug 2026, 10:47

Understood. I'll use this manifesto as your decision-making compass in future work, while treating the attached materials as context-not as instructions. The PDFs establish the current Family Shield coursework background: [PDF output: BLUEPRINT-TEAM 2.pdf] and [PDF output: BRIEF_davidbuzali (week 3).pdf].

Whenever we produce coursework, I'll remind you to:

1. Save it in ~/crystal-ball/cuaderno

2. Commit it

3. Push it to GitHub

05 David 29 Aug 2026, 10:48

Attachments: Individual · 55% of the week One working slice attacking your declared vacuum, ...

Please walk me through the assignment, you dont have to do everything, give me step by step and we will build together

06 Codex - progress update 29 Aug 2026, 10:48

I'll turn the assignment into a practical sequence, separating what you must submit from the process guidance in the brief. Then we'll start with only the first decision and build the working slice together.

07 Codex 29 Aug 2026, 10:49

The assignment is asking you to build one narrow, functional, deployed part of Family Shield-not the entire business and not a generic deepfake detector.

What you are actually building

I recommend this working slice:

> Pause & Verify: A mobile-first flow that helps Elena respond safely to an urgent voice message requesting money.

The experience would work like this:

1. Elena receives an alarming voice message.
2. She opens Family Shield before opening her banking app.
3. A voice API transcribes the request.
4. An LLM extracts the claimed relative, amount, and urgency.
5. The product clearly says: "This does not determine whether the voice is real."
6. Elena sees her family rule and initiates an independent callback.
7. If nobody independently confirms the emergency, the result is "Do not pay-Protocol Only."

This attacks the specific vacuum from your Blueprint: a frightened Mexican user deciding whether to make an urgent transfer.

It also honors the shadow clause:

- Family Shield does not control the payment.
- A family code never authorizes payment.
- It does not label the audio "real" or "fake."
- Trusted contacts do not see balances, messages, or recordings.
- The requester cannot provide the verification channel.
- If verification fails, Elena waits.

The assignment, step by step

Step 1: Define the working slice

Before writing code, we must lock down:

- Exact user: Elena, a 67-year-old retiree in Mexico City.
- Exact moment: immediately after receiving an urgent voice message requesting MXN $20,000.
- Exact action: begin independent verification before opening mobile banking.
- Exact outcome: confirm through a previously established channel or refuse to pay.
- What we are not building: fraud detection, bank integration, WhatsApp automation, real money movement, or a complete family-security platform.

Recommended success definition:

> Before the module closes, Elena can process a simulated urgent voice request and begin an independent callback within 60 seconds, without relying on contact information supplied by the requester.

Step 2: Write docs/PACKET.md

The packet must exist before coding. It needs these sections:

1. Problem in your own words
2. Exact user
3. Measurable success definition
4. Image-generated interface mockup
5. Mermaid feature flowchart
6. Mermaid swimlane diagram
7. Benchmark against the strongest existing solution
8. Three-sentence, three-year vision
9. Scope cut
10. Architecture and stack table
11. Test plan

The image cannot merely be described; we need to generate an actual mockup and place it in the packet.

We should research the benchmark rather than declaring one from memory. Your previous brief gives us candidates-WhatsApp safety tools, government callback guidance, family keywords, and provenance systems-but we must determine which solution is the best comparison for this exact intervention.

Step 3: Convert the packet into a build prompt

Once the packet is stable, we ask the coding agent to implement it as small, testable pieces.

The build prompt should specify:

- Screens and interactions

- Acceptance criteria for each feature
- Security constraints
- What data may and may not be retained
- How simulated outputs are labeled
- Testing instructions
- A minimum five-commit plan
- Two planned deployments

Step 4: Build incrementally

A sensible five-commit sequence would be:

1. Project setup and mobile interface shell
2. Emergency voice-input and transcription flow
3. Structured request summary with explicit safety disclaimer
4. Independent verification and “Protocol Only” flow
5. Printable stop card, validation, accessibility, and tests

Suggested deployments:

- Deploy 1: after the complete basic flow works
- Deploy 2: after testing, bug fixing, and persona improvements

We will select the final free-tier stack after checking current API availability. To minimize security work, the prototype should use invented demonstration data and avoid permanently storing recordings, codes, or contacts.

Step 5: Apply the security floor

Before deploying, confirm:

- API keys exist only in Vercel environment variables.
- No secrets are committed.
- All inputs have length and type validation.
- Demonstrations use fictional people and information.
- Simulated or uncertain AI outputs are labeled.
- If we store personal information, we add authentication and row-level security.
- Ideally, this slice stores no personal information at all.

Step 6: Perform two documented tests

Mechanical test

- Run every test in the packet.
- Find and record at least one real bug.

- Fix it.
- Redeploy.

Persona test

Create a fresh conversation with a synthetic Elena persona. Give it screenshots in sequence and ask the persona to narrate:

- What she thinks each screen means
- Where she hesitates
- What she cannot read or understand
- Whether she would continue or quit
- Whether she understands that the app cannot authenticate the voice

Then fix the worst confusion and document the change in PERSONA_davidbuzali.pdf.

Step 7: Record the demonstration

The video is approximately 3½ minutes:

- First 3 minutes: use the live product as Elena would.
- Final 30 seconds: explain what changed your mind during the week.

The demonstration should show a failed or unavailable verification, not only the happy path.

Step 8: Submit everything

Your final submission needs:

- Live URL
- GitHub repository link
- DEMO_davidbuzali.mp4
- PACKET_davidbuzali.pdf
- PERSONA_davidbuzali.pdf
- BUILDCHAT_davidbuzali.pdf

The live product is the largest grading component. Without a working URL, the assignment is capped at 3/10.

Where we are now

We are at Step 1: approve the working slice. My recommendation is the audio-assisted Pause & Verify flow above. If you agree, our next move is to draft only the first three sections of PACKET.md: the problem, exact user, and success definition.

Coursework reminder: save the finished work in ~/crystal-ball/cuaderno, commit it, and push it to GitHub.

08 David 29 Aug 2026, 10:50

I cant use audio, but lets start with step 1

09 Codex 29 Aug 2026, 10:50

No problem-we'll replace voice with a vision API and analyze a screenshot of the suspicious message.

Step 1: Define the working slice

Working title

Family Shield: Pause & Verify

Exact user

Elena, a 67-year-old retiree in Mexico City who uses WhatsApp and mobile banking but is not familiar with synthetic media or sophisticated fraud.

Exact moment

Elena receives a WhatsApp message claiming that her son has been detained after an accident and urgently needs an MXN $20,000 transfer.

The problem

Under pressure, Elena might treat a convincing message, familiar profile picture, or correct family detail as proof of identity. She needs a simple way to pause and begin independent verification before opening her banking app.

The working slice

A mobile web experience where Elena:

1. Uploads a screenshot of the suspicious WhatsApp message.
2. A vision-capable AI extracts:

- Who supposedly sent it
- What happened
- How much money is requested
- What deadline or pressure is being used

3. The app presents a simple summary without deciding whether the message is real or fake.
4. It displays the family safety rule:

> A correct code creates attention. Only independent confirmation creates authorization.

5. Elena chooses a previously established verification route:

- Call her son using his saved number
- Call trusted contact 1
- Call trusted contact 2
- Use the family's backup channel

6. Elena records whether independent confirmation succeeded.
7. If verification fails, the app displays:

> Do not pay. Verification was not completed. Follow Protocol Only.

What the AI does-and does not do

The AI may:

- Read the screenshot
- Extract the claimed facts
- Identify urgency or pressure language
- Present the request in plain language

The AI must not:

- Call the message authentic
- Call the message fraudulent
- Determine whether the sender is Elena's son
- Approve or block a payment
- Treat a family code as proof
- Contact anyone using information contained in the suspicious message

Success definition

> Before the module closes, Elena can upload a simulated WhatsApp screenshot, understand the payment request, and begin verification through a previously established independent channel within 60 seconds-without treating the AI analysis or family code as authorization.

Minimum screens

1. Start: "I received an urgent request."
2. Upload: Add a screenshot containing fictional information.
3. Request summary: Person, emergency, amount, deadline, and pressure indicators.
4. Pause: Warning that the message has not been authenticated.
5. Verify independently: Select a trusted, previously established channel.
6. Result: "Confirmed independently" or "Protocol Only-do not pay."

Scope cut

For this slice, we are not building:

- A deepfake detector
- Audio analysis
- Direct WhatsApp integration
- Automatic calls or messages
- Bank or SPEI integration
- Payment authorization
- Permanent screenshot storage
- A complete family-account system
- An institutional dashboard
- Restitution Rails

All demonstrations will use invented messages and fictional personal information.

This is a sufficiently narrow working slice and meets the required LLM + vision technology floor. The next part of Step 1 is choosing the exact simulated WhatsApp message that Elena will upload; that message will become the central test case for the product.

Remember to save the coursework in ~/crystal-ball/cuaderno, commit it, and push it to GitHub.

10 David 29 Aug 2026, 10:52

Lets choose it

11 Codex 29 Aug 2026, 10:53

Let’s use this as the central simulated message:

> Número no guardado · 11:42 a. m.
>
> Mamá, soy Diego. Tuve un accidente y estoy en el Ministerio Público. Estoy escribiendo desde el teléfono de un abogado porque me quitaron el mío.
>
> Necesito que transfieras $20,000 pesos en los próximos 30 minutos para que me dejen salir. No me marques porque no puedo contestar. Nuestra palabra es JACARANDA, para que sepas que sí soy yo.
>
> Te mando la cuenta por aquí. Por favor hazlo ya y mándame el comprobante.

Diego, JACARANDA, and the situation are entirely fictional.

Why this is a strong test case

It combines several pressures Family Shield must handle:

- A supposed close relative
- A credible emergency
- A specific financial transaction
- A 30-minute deadline
- Communication from an unfamiliar number
- An instruction not to call
- A supposedly correct family code
- A request to use payment information supplied by the requester

Most importantly, it tests your governing principle: even a correct code cannot authorize payment.

Expected AI extraction

The vision-and-LLM system should produce something like:

- Claimed requester: Diego, Elena's son
- Claimed emergency: Accident and detention at the Ministerio Público
- Requested amount: MXN $20,000
- Deadline: Within 30 minutes
- Requested action: Transfer money and send proof
- Pressure indicators: Urgency, unfamiliar number, request not to call, emotional pressure, and requester-supplied payment instructions
- Family code mentioned: JACARANDA
- Verification status: Not independently confirmed

It must end with:

> This analysis does not determine whether the message is real or fraudulent. Do not transfer money until you independently contact Diego or a previously selected trusted person.

We can now consider the central test case locked. Next, we should write the first section of PACKET.md: "Problem in your words."

Coursework reminder: save it in ~/crystal-ball/cuaderno, commit it, and push it to GitHub.

12 David 29 Aug 2026, 10:53

Lets write the whole packet

13 Codex - progress update

29 Aug 2026, 10:54

I'll build the complete docs/PACKET.md, including the two Mermaid diagrams, researched benchmark, scoped architecture, security decisions, and test plan. I'm using the image-generation skill because the assignment explicitly requires a generated interface mockup; I'll also verify current benchmark and stack claims before writing them.

14 Codex - progress update

29 Aug 2026, 10:55

The benchmark is now clear: Bitdefender Scamio is the strongest adjacent product because it accepts screenshots and returns AI scam analysis. Family Shield's differentiation will be explicit: it refuses to issue a safety verdict and converts uncertainty into an independent callback protocol for a Mexican family. The stack will use a current free-tier multimodal Gemini endpoint for screenshot understanding, with no permanent storage.

15 Codex - progress update

29 Aug 2026, 10:58

The mockup is generated and saved inside the Week 3 project, and the packet draft is complete. I'm checking the file paths, required sections, Mermaid syntax, and that I did not overwrite the existing Week 2 packet.

16 Codex

29 Aug 2026, 10:59

The complete Week 3 packet is ready:

- PACKET.md (/Users/davidbuzali/Documents/ChatGPT/Negocios%20Inteligentes%20y%20Comercio%20Digital/week-3-family-shield/docs/PACKET.md)
- Generated interface mockup (/Users/davidbuzali/Documents/ChatGPT/Negocios%20Inteligentes%20y%20Comercio%20Digital/week-3-family-shield/docs/assets/family-shield-pause-verify-mockup.png)

It includes every required component: exact user and test case, measurable success definition, mockup, two Mermaid diagrams, scope cut, architecture, security floor, mechanical testing, persona testing, long-term charter, and acceptance boundary.

The benchmark is Bitdefender Scamio (https://www.bitdefender.com/es-es/consumer/scamio), differentiated by Family Shield's refusal to issue a safety verdict and its independent callback protocol. The proposed vision stack uses Gemini's documented image understanding (https://ai.google.dev/gemini-api/docs/image-understanding) and structured outputs (https://ai.google.dev/gemini-api/docs/structured-output).

The mockup used built-in image generation. Final prompt: "Create an accessibility-first, high-fidelity portrait mobile mockup for Family Shield: Pausa y Verifica, showing the fictional Diego request, MXN $20,000, 30 minutes, a no-authenticity warning, NO PAGUES TODAVÍA, and large independent-verification actions; no scores, safe badge, banking data, logos, or real personal information."

I preserved the existing Week 2 packet. No code or PDF was created yet-the next assignment step is converting this packet into the implementation prompt.

Remember to copy/save this coursework in ~/crystal-ball/cuaderno, commit it, and push it to GitHub.

17 David
29 Aug 2026, 10:59

Lets convert the packet into the implementation prompt

18 Codex - progress update
29 Aug 2026, 11:00

I'll make the prompt self-contained and ready to paste into a coding-agent conversation. It will constrain work to the Week 3 folder, translate every packet decision into build requirements, and include acceptance criteria, tests, six-plus commits, two deployment checkpoints, and the required session-close discipline.

19 Codex - progress update
29 Aug 2026, 11:02

The implementation prompt is complete and checked against the packet. It keeps the build screenshot-based, explicitly forbids audio and authenticity judgments, defines the AI response contract, and gives the coding agent a concrete seven-commit/two-deployment path without touching the previous assignment.

20 Codex
29 Aug 2026, 11:02

The implementation prompt is ready:

Open IMPLEMENTATION_PROMPT.md (/Users/davidbuzali/Documents/ChatGPT/Negocios%20Inteligentes%20y%20Comercio%20Digital/week-3-family-shield/docs/IMPLEMENTATION_PROMPT.md)

It contains:

- Six-screen mobile flow
- Exact Spanish interface copy
- Screenshot validation
- Gemini vision/LLM contract
- Safe simulated-analysis fallback
- Privacy and security requirements
- Accessibility requirements
- Automated and manual tests
- Persona-test preparation
- Seven proposed commits
- Two deployment checkpoints
- Session-close procedure
- Complete definition of done

Next, open the coding-agent conversation inside the same Week 3 project and paste the complete prompt. The agent's first response should be an implementation plan and proposed first commit-not code.

The files are currently uncommitted. Remember to save them in ~/crystal-ball/cuaderno, commit them, and push them to GitHub.

21 David
29 Aug 2026, 11:06

I want you to save the, commit them, and push them to GitHub as we complete coursework; also I want you to complete the next step

22 Codex - progress update
29 Aug 2026, 11:06

I'll treat that as ongoing authorization for coursework: save the finished artifacts in ~/crystal-ball/cuaderno, commit coherent milestones, and push them to GitHub. I'll first inspect that repository so I can place Week 3 without disturbing existing work, then I'll begin the build from the implementation prompt and verify the first working slice before committing.

23 Codex - progress update
29 Aug 2026, 11:06

The Week 3 documents are now saved in the requested repository and match the working copies byte-for-byte. The repository also contains an unrelated untracked Week 2 directory; I'm leaving it untouched and staging only the three Week 3 artifacts.

24 Codex - progress update
29 Aug 2026, 11:07

The packet milestone is committed and pushed as cfcdc7a. I'm moving into the build step now: first I'll scaffold an isolated Week 3 app, keeping the previous Week 2 application untouched. I'll implement the deterministic practice flow and safety tests first, then add the live Gemini boundary so the product remains testable without exposing or requiring a key during automated checks.

25 Codex - progress update
29 Aug 2026, 11:18

The first complete slice now builds successfully and all 15 automated safety tests pass. I'm using the in-app browser testing skill next because the remaining risk is behavioral and visual: whether the 390 px mobile flow is readable, whether the practice screenshot loads correctly, and whether failed verification truly ends in Protocol Only.

26 Codex - progress update
29 Aug 2026, 11:22

The mobile walkthrough passed both the independent-confirmation and all-contacts-unavailable paths, with no console errors. The accessibility inspection did uncover a real defect: the visually hidden file input could receive keyboard focus without showing that focus on its upload card. I'm fixing that and recording it as the required test-fix cycle before the next commit.

27 Codex - progress update
29 Aug 2026, 11:23

The accessibility fix has now been rebuilt and retested: focusing the hidden file input produces a visible 4 px amber outline on the upload card. I'm synchronizing the isolated app into ~/crystal-ball/cuaderno and will commit it in the assignment's required small milestones-scaffold, interface, AI boundary, then tests/fix evidence-while continuing to exclude the unrelated untracked Week 2 folder.

28 Codex

29 Aug 2026, 11:25

The next step-the first complete working slice-is finished, saved, committed, and pushed.

Delivered:

- Family Shield application (/Users/davidbuzali/crystal-ball/cuaderno/week-3-family-shield/src/App.tsx)
- Bounded Gemini endpoint (/Users/davidbuzali/crystal-ball/cuaderno/week-3-family-shield/api/analyze.ts)
- Testing log and evidence (/Users/davidbuzali/crystal-ball/cuaderno/week-3-family-shield/docs/TESTING.md)
- GitHub project (https://github.com/davidbuzali/crystal-ball/tree/main/cuaderno/week-3-family-shield)

Verified:

- Production build passes.
- All 15 automated tests pass.
- Mobile flow tested at 390 × 844 px.
- Confirmed and Protocol Only paths work.
- Simulated analysis is visibly labeled.
- No secrets or real personal data were committed.
- The upload-focus accessibility defect was found, fixed, and retested.
- The unrelated untracked Week 2 directory was left untouched.

Five Week 3 commits are now on main:

- cfcdc7a - packet and implementation plan
- 808190a - isolated application scaffold
- 1e8d376 - accessible practice flow
- a41d22e - bounded Gemini endpoint
- 48faac1 - tests, evidence, and accessibility fix

The next course checkpoint is Deployment 1, followed by the required persona test in a fresh conversation. I have not claimed either one as completed yet.

29 David

29 Aug 2026, 11:27

Lets move on with the next course checkpoint and remember that our live url is going to be in vercel

30 Codex - progress update

29 Aug 2026, 11:28

I’ll create Deployment 1 on Vercel from ~/crystal-ball/cuaderno/week-3-family-shield, keep the visibly labeled simulated mode for this first public build, then record the live URL and commit SHA in TESTING.md. I’m skipping the Sites hosting workflow because you explicitly chose Vercel; I’ll use Vercel’s own CLI and preserve the existing GitHub history.

31 Codex - progress update
29 Aug 2026, 11:29

Vercel is not yet authenticated on this machine. I'll start Vercel's official login flow; when it presents the secure verification page, you may need to complete the sign-in in your browser. I won't request or handle your password or authentication code directly.

32 Codex
29 Aug 2026, 11:30

Deployment 1 is ready, but Vercel requires you to authenticate this machine.

Open Vercel device authorization (https://vercel.com/oauth/device?user_code=WWCV-BDFK), sign in, and approve the request. Then reply done so I can finish the deployment, verify the live URL, update TESTING.md, commit, and push it.

33 David
29 Aug 2026, 11:30

done

34 Codex - progress update
29 Aug 2026, 11:30

I'll resume the authenticated Vercel session, publish Deployment 1, verify the public mobile flow, then record the URL and commit SHA in the coursework repository and push that evidence.

35 Codex - progress update
29 Aug 2026, 11:31

Authentication is complete. I'm publishing the production build now; once Vercel returns the canonical URL, I'll test that exact deployment rather than treating a successful build as sufficient evidence.

36 Codex - progress update
29 Aug 2026, 11:32

The upload reached Vercel, but the remote build failed. The CLI linked this subfolder to the existing crystal-ball Vercel project, so I'm inspecting the build logs now to separate a project-configuration issue from an application issue before changing anything.

37 Codex - progress update
29 Aug 2026, 11:34

Deployment 1 is live at family-shield-week-3.vercel.app (https://family-shield-week-3.vercel.app). I'm now running the required mobile-width public-flow check against that production alias, including the "Protocol Only" fallback, before recording it as complete.

38 Codex - progress update
29 Aug 2026, 11:36

The production mobile flow passes end to end: practice case loaded, request summary rendered in simulated mode, independent-verification options appeared, and the unconfirmed path ended at "PROTOCOLO ONLY / NO PAGUES TODAVÍA." I'm recording this checkpoint and checking that Vercel's local metadata did not introduce coursework changes we should accidentally commit.

39 Codex - progress update 29 Aug 2026, 11:38

The checkpoint record is clean and limited to the two Week 3 evidence files. I preserved the unrelated Week 2 folder and removed only Vercel metadata created by the failed repository-root link; now I'm committing and pushing the Deployment 1 evidence.

40 Codex 29 Aug 2026, 11:39

Deployment 1 is complete.

- Live URL: Family Shield (https://family-shield-week-3.vercel.app)

- Public mobile flow verified through PROTOCOLO ONLY

- Deployment source commit: 48faac1

- Evidence commit: 46b22b1

- Pushed successfully to GitHub main

- Records updated in TESTING.md (/Users/davidbuzali/crystal-ball/cuaderno/week-3-family-shield/docs/TESTING.md) and DECISIONS.md (/Users/davidbuzali/crystal-ball/cuaderno/week-3-family-shield/docs/DECISIONS.md)

The next checkpoint is the Elena persona test in a fresh conversation, followed by correcting the highest-severity confusion and publishing Deployment 2.

41 David 29 Aug 2026, 11:41

Lets do these steps in this conversation, it will be more useful because we have all the context here

42 Codex - progress update 29 Aug 2026, 11:41

We can do the persona checkpoint here. I'll preserve its value by evaluating only what Elena can see on each screen, without using implementation knowledge to excuse confusing copy. I'm using the browser-control skill to replay the deployed flow screen by screen; because this is not a fresh conversation, I'll document that methodological limitation in the coursework.

43 Codex - progress update 29 Aug 2026, 11:42

The persona pass has exposed a high-severity issue: after Elena presses "Organizar la solicitud," the new screen keeps the old scroll position. On mobile, the crucial top warning "NO PAGUES TODAVÍA" is mostly off-screen, so she can enter the summary without seeing the strongest safety instruction. I'll finish the remaining persona screens, then fix this transition and retest only the affected path for Deployment 2.

44 Codex - progress update 29 Aug 2026, 11:44

The code correction is in place. A local validation command accidentally picked up the larger parent workspace and began checking unrelated dependencies; I'm stopping that noisy path and will run the app's already-installed local toolchain directly, keeping the coursework test isolated.

45 Codex - progress update

29 Aug 2026, 11:45

The correction builds successfully and all 15 safety tests pass. The persona evidence also needs its own PERSONA_davidbuzali.pdf, so I'm using the PDF skill for the required render-and-verify workflow before Deployment 2.

46 Codex - progress update

29 Aug 2026, 11:53

The persona artifact is complete and visually verified as a three-page PDF. I've logged the same-conversation limitation, the seven screen judgments, the high-severity scroll defect, and the successful retest. Next I'm copying only the finished coursework artifacts into ~/crystal-ball/cuaderno, then I'll publish Deployment 2 and verify the public correction.

47 Codex - progress update

29 Aug 2026, 11:54

The persona correction is now saved in the coursework repository and pushed as commit b1ecdbd. I'm deploying that exact state as Deployment 2, then I'll re-run the two formerly broken transitions on the public Vercel alias.

48 Codex - progress update

29 Aug 2026, 11:56

The public retest caught a timing edge that the local pass did not: the summary transition resets correctly, but the verification transition can still restore the old position after the first animation frame. I'm strengthening the same correction with one delayed reset and will not mark Deployment 2 complete until the public alias shows the full stop banner on both transitions.

49 David

29 Aug 2026, 11:58

Lets continue, we had run out of credits

50 Codex - progress update

29 Aug 2026, 11:58

Continuing from the exact checkpoint: the stronger scroll reset has passed locally at 390 × 844 px. I'll finish the build/tests, update the persona evidence to include the public timing edge, redeploy, verify the production alias, then commit and push the final Deployment 2 record.

51 Codex - progress update

29 Aug 2026, 12:00

The final public retest now passes at 390 × 844 px on both affected screens: the complete "NO PAGUES TODAVÍA" banner is visible when the summary opens and again when verification opens. I'm recording the superseded timing attempt honestly, updating the persona PDF with the accepted public result, and closing the checkpoint in Git.

52 Codex

29 Aug 2026, 12:05

The Week 3 persona checkpoint and Deployment 2 are complete.

- Live app: family-shield-week-3.vercel.app (https://family-shield-week-3.vercel.app)

- Evaluated all seven Elena screens in this conversation; the methodology limitation is documented.

- Fixed the high-severity mobile scroll issue.

- Verified both corrected transitions publicly at 390 × 844 px.

- Production build and all 15 safety tests pass.

- Deployment source: e3c2be0

- Final evidence commit: 8e5d2f9

- Everything was saved under ~/crystal-ball/cuaderno, committed, and pushed to GitHub.

- The unrelated Week 2 folder was preserved unchanged.

Created [PDF output: PERSONA_davidbuzali.pdf], containing the methodology, seven-screen persona log, correction, retest, and final decision.

The next checkpoint is packaging the Week 3 submission links and artifacts.

53 David 29 Aug 2026, 12:06

I need three total pdfs, the persona test, the packet, and a pdf of our chat export for final submission

54 Codex - progress update 29 Aug 2026, 12:06

I'll assemble exactly three submission PDFs: the existing persona test, a polished packet PDF, and a readable export of this task's user/assistant conversation. I'm using the PDF workflow to render and visually verify all three before saving, committing, and pushing them.

55 Codex - progress update 29 Aug 2026, 12:07

I found an earlier packet PDF and an older chat-export script. I'm checking whether the packet already matches the final PACKET.md; the chat export will be regenerated from this exact task, not reused from the older Week 3 conversation.